

```
1 !wget https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
2 !unzip jena_climate_2009_2016.csv.zip
```

```
--2024-11-02 17:56:08-- https://s3.amazonaws.com/keras-datasets/jena_climate_2009_2016.csv.zip
Resolving s3.amazonaws.com (s3.amazonaws.com)... 16.15.193.111, 52.217.126.56, 54.231.194.128, ...
Connecting to s3.amazonaws.com (s3.amazonaws.com)|16.15.193.111|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 13565642 (13M) [application/zip]
Saving to: 'jena_climate_2009_2016.csv.zip.1'
```

```
jena_climate_2009_2 100%[=====>] 12.94M 15.9MB/s in 0.8s
```

```
2024-11-02 17:56:10 (15.9 MB/s) - 'jena_climate_2009_2016.csv.zip.1' saved [13565642/13565642]
```

```
Archive: jena_climate_2009_2016.csv.zip
replace jena_climate_2009_2016.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename:
```

```
1 import os
2 data = os.path.join("jena_climate_2009_2016.csv")
```

```
1 with open(data) as f:
2     data = f.read()
```

```
1 records = data.split("\n")
2 title = records[0].split(",")
3 records = records[1:]
4 print(title)
5 print(len(records))
6 import os
7 data = os.path.join("jena_climate_2009_2016.csv")
```

```
['Date Time', 'p (mbar)', 'T (degC)', 'Tpot (K)', 'Tdew (degC)', 'rh (%)', 'VPmax (mbar)', 'VPact (mbar)'
420451
```

```
1 with open(data) as f:
2     data = f.read()
```


```
1 records = data.split("\n")
2 title = records[0].split(",")
3 records = records[1:]
4 print(title)
5 print(len(records))
6 #Processing and analyzing the data.
7
```

```
['Date Time', 'p (mbar)', 'T (degC)', 'Tpot (K)', 'Tdew (degC)', 'rh (%)', 'VPmax (mbar)', 'VPact (mbar)'
420451
```

```
1 import numpy as np
2 temp = np.zeros((len(records),))
3 original_data = np.zeros((len(records), len(title) - 1))
4 for i, line in enumerate(records):
5     values = [float(x) for x in line.split(",")[1:]]
6     temp[i] = values[1]
7     original_data[i, :] = values[:]
8 import numpy as np
9 temp = np.zeros((len(records),))
10 original_data = np.zeros((len(records), len(title) - 1))
11 for i, line in enumerate(records):
12     values = [float(x) for x in line.split(",")[1:]]
13     temp[i] = values[1]
14     original_data[i, :] = values[:]
15 #Visualizing the temperature time series data.
```


```
1 from matplotlib import pyplot as plt
2 plt.plot(range(len(temp)), temp)
3 #Visualizing the temperature time series for the first 10 days.
```

↗ [`<matplotlib.lines.Line2D at 0x7f7bbd69b490>`]



```
1 plt.plot(range(1440), temp[:1440])
2 #Calculating the number of samples for each data split.
```

↗ [`<matplotlib.lines.Line2D at 0x7f7bf41ea4a0>`]



```
1 train_Sample_Num = int(0.5 * len(original_data))
2 val_original_data = int(0.25 * len(original_data))
3 test_Sample_Number = len(original_data) - train_Sample_Num - val_original_data
4 print("train_Sample_Num:", train_Sample_Num)
5 print("val_original_data:", val_original_data)
6 print("test_Sample_Number:", test_Sample_Number)
7 #Data Preparation
8 #Normalizing the data
```

↗

```
train_Sample_Num: 210225
val_original_data: 105112
test_Sample_Number: 105114
```

```
1 mean = original_data[:train_Sample_Num].mean(axis=0)
2 original_data -= mean
3 std = original_data[:train_Sample_Num].std(axis=0)
4 original_data /= std
5 import numpy as np
6 from tensorflow import keras
7 num_series = np.arange(10)
8 dummy_dataset = keras.utils.timeseries_dataset_from_array(
9     data=num_series[:3],
10     targets=num_series[3:],
11     sequence_length=3,
12     batch_size=2,
13 )
```

```
1 for inputs, targets in dummy_dataset:
2     for i in range(inputs.shape[0]):
```

```

3         print([int(x) for x in inputs[i]], int(targets[i]))
4 #Instantiating datasets for training, validation, and testing

↪ [0, 1, 2] 3
   [1, 2, 3] 4
   [2, 3, 4] 5
   [3, 4, 5] 6
   [4, 5, 6] 7

1 sample_rate = 6
2 seq_len = 120
3 delay = sample_rate * (seq_len + 24 - 1)
4 batch_size = 256

1 train_dataset = keras.utils.timeseries_dataset_from_array(
2     original_data[:~delay],
3     targets=temp[~delay:],
4     sampling_rate=sample_rate,
5     sequence_length=seq_len,
6     shuffle=True,
7     batch_size=batch_size,
8     start_index=0,
9     end_index=train_Sample_Num)

1 val_dataset = keras.utils.timeseries_dataset_from_array(
2     original_data[:~delay],
3     targets=temp[~delay:],
4     sampling_rate=sample_rate,
5     sequence_length=seq_len,
6     shuffle=True,
7     batch_size=batch_size,
8     start_index=train_Sample_Num,
9     end_index=train_Sample_Num + val_original_data)

1 test_dataset = keras.utils.timeseries_dataset_from_array(
2     original_data[:~delay],
3     targets=temp[~delay:],
4     sampling_rate=sample_rate,
5     sequence_length=seq_len,
6     shuffle=True,
7     batch_size=batch_size,
8     start_index=train_Sample_Num + val_original_data)
9 #Examining the results of one of our datasets to assess its content and quality.

1 for samples, targets in train_dataset:
2     print("samples shape:", samples.shape)
3     print("targets shape:", targets.shape)
4     break
5 #A common-sense, non-machine-learning baseline
6 #Computing the common-sense baseline MAE

↪ samples shape: (256, 120, 14)
   targets shape: (256,)

1 def evaluate_naive_method(dataset):
2     total_abs_err = 0.
3     samples_seen = 0
4     for samples, targets in dataset:
5         preds = samples[:, -1, 1] * std[1] + mean[1]
6         total_abs_err += np.sum(np.abs(preds - targets))
7         samples_seen += samples.shape[0]
8     return total_abs_err / samples_seen

1 print(f"Validation MAE: {evaluate_naive_method(val_dataset):.2f}")
2 print(f"Test MAE: {evaluate_naive_method(test_dataset):.2f}")
3

↪ Validation MAE: 2.44
   Test MAE: 2.62

1 !pip install tensorflow==2.12

```

```

Collecting tensorflow==2.12
  Downloading tensorflow-2.12.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (3.4 kB)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (1.6)
Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (24.3)
Collecting gast<=0.4.0,>=0.2.1 (from tensorflow==2.12)
  Downloading gast-0.4.0-py3-none-any.whl.metadata (1.1 kB)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (1)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (3.12.1)
Requirement already satisfied: jax>=0.3.15 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (0.4.33)
Collecting keras<2.13,>=2.12.0 (from tensorflow==2.12)
  Downloading keras-2.12.0-py2.py3-none-any.whl.metadata (1.4 kB)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (18.1)
Collecting numpy<1.24,>=1.22 (from tensorflow==2.12)
  Downloading numpy-1.23.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.3 kB)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (3.4)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (24.1)
Requirement already satisfied: protobuf!=4.21.0,!<4.21.1,!<4.21.2,!<4.21.3,!<4.21.4,!<4.21.5,<5.0.0dev,>=3.20.3 in /usr/
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (75.1.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (1.16.0)
Collecting tensorboard<2.13,>=2.12 (from tensorflow==2.12)
  Downloading tensorboard-2.12.3-py3-none-any.whl.metadata (1.8 kB)
Collecting tensorflow-estimator<2.13,>=2.12.0 (from tensorflow==2.12)
  Downloading tensorflow_estimator-2.12.0-py2.py3-none-any.whl.metadata (1.3 kB)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12) (2.5)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12)
Collecting wrapt<1.15,>=1.11.0 (from tensorflow==2.12)
  Downloading wrapt-1.14.1-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12)
Requirement already satisfied: jaxlib<=0.4.33,>=0.4.33 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12)
Requirement already satisfied: ml-dtypes>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12)
INFO: pip is looking at multiple versions of jax to determine which version is compatible with other requirements. This
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.35-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.35,>=0.4.34 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.35-cp310-cp310-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.34-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.34,>=0.4.34 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.34-cp310-cp310-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.31-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.31,>=0.4.30 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.31-cp310-cp310-manylinux2014_x86_64.whl.metadata (983 bytes)
Collecting jax>=0.3.15 (from tensorflow==2.12)
  Downloading jax-0.4.30-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.30,>=0.4.27 (from jax>=0.3.15->tensorflow==2.12)
  Downloading jaxlib-0.4.30-cp310-cp310-manylinux2014_x86_64.whl.metadata (1.0 kB)
Requirement already satisfied: scipy>=1.9 in /usr/local/lib/python3.10/dist-packages (from jax>=0.3.15->tensorflow==2.12)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.12)
Collecting google-auth-oauthlib<1.1,>=0.5 (from tensorboard<2.13,>=2.12->tensorflow==2.12)
  Downloading google_auth_oauthlib-1.0.0-py2.py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.13,>=2.12-)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.13,>=2.12-)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from te
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.13,>=2.12-)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->ten
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->ten
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->ten
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from google-auth-oau
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tenso
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tenso
Requirement already satisfied: pyasn1<0.7.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages (from pyasn1-modules>=0.2
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from requests-oauthlib>=0.7.0)
Download tensorflow-2.12.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (585.9 MB)
585.9/585.9 MB 2.8 MB/s eta 0:00:00
Download gast-0.4.0-py3-none-any.whl (9.8 kB)
Download jax-0.4.30-py3-none-any.whl (2.0 MB)
2.0/2.0 MB 79.8 MB/s eta 0:00:00
Download keras-2.12.0-py2.py3-none-any.whl (1.7 MB)
1.7/1.7 MB 77.4 MB/s eta 0:00:00
Download numpy-1.23.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.1 MB)
17.1/17.1 MB 77.3 MB/s eta 0:00:00
Download tensorboard-2.12.3-py3-none-any.whl (5.6 MB)
5.6/5.6 MB 106.3 MB/s eta 0:00:00
Download tensorflow_estimator-2.12.0-py2.py3-none-any.whl (440 kB)
440.7/440.7 kB 33.8 MB/s eta 0:00:00
Download wrapt-1.14.1-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2014_x86_64.w
77.9/77.9 kB 7.5 MB/s eta 0:00:00
Download google_auth_oauthlib-1.0.0-py2.py3-none-any.whl (18 kB)
Download jaxlib-0.4.30-cp310-cp310-manylinux2014_x86_64.whl (79.6 MB)
79.6/79.6 MB 9.8 MB/s eta 0:00:00
Installing collected packages: wrapt, tensorflow-estimator, numpy, keras, gast, jaxlib, google-auth-oauthlib, tensorboar
Attempting uninstall: wrapt
Found existing installation: wrapt 1.16.0

```

```
Uninstalling wrapt-1.16.0:
Successfully uninstalled wrapt-1.16.0
Attempting uninstall: numpy
Found existing installation: numpy 1.26.4
Uninstalling numpy-1.26.4:
Successfully uninstalled numpy-1.26.4
Attempting uninstall: keras
Found existing installation: keras 3.4.1
Uninstalling keras-3.4.1:
Successfully uninstalled keras-3.4.1
Attempting uninstall: gast
Found existing installation: gast 0.6.0
Uninstalling gast-0.6.0:
Successfully uninstalled gast-0.6.0
Attempting uninstall: jaxlib
Found existing installation: jaxlib 0.4.33
Uninstalling jaxlib-0.4.33:
Successfully uninstalled jaxlib-0.4.33
Attempting uninstall: google-auth-oauthlib
Found existing installation: google-auth-oauthlib 1.2.1
Uninstalling google-auth-oauthlib-1.2.1:
Successfully uninstalled google-auth-oauthlib-1.2.1
Attempting uninstall: tensorboard
Found existing installation: tensorboard 2.17.0
Uninstalling tensorboard-2.17.0:
Successfully uninstalled tensorboard-2.17.0
Attempting uninstall: jax
Found existing installation: jax 0.4.33
Uninstalling jax-0.4.33:
Successfully uninstalled jax-0.4.33
Attempting uninstall: tensorflow
Found existing installation: tensorflow 2.17.0
Uninstalling tensorflow-2.17.0:
Successfully uninstalled tensorflow-2.17.0
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour
albucore 0.0.19 requires numpy>=1.24.4, but you have numpy 1.23.5 which is incompatible.
alumentations 1.4.20 requires numpy>=1.24.4, but you have numpy 1.23.5 which is incompatible.
bigframes 1.25.0 requires numpy>=1.24.0, but you have numpy 1.23.5 which is incompatible.
chex 0.1.87 requires numpy>=1.24.1, but you have numpy 1.23.5 which is incompatible.
tf-keras 2.17.0 requires tensorflow<2.18,>=2.17, but you have tensorflow 2.12.0 which is incompatible.
xarray 2024.10.0 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
Successfully installed gast-0.4.0 google-auth-oauthlib-1.0.0 jax-0.4.30 jaxlib-0.4.30 keras-2.12.0 numpy-1.23.5 tensorbo
```

```

1 from tensorflow import keras
2 from tensorflow.keras import layers
3
4 inputs = keras.Input(shape=(seq_len, original_data.shape[-1]))
5 x = layers.Flatten()(inputs)
6 x = layers.Dense(64, activation="relu")(x)
7 outputs = layers.Dense(1)(x)
8 model = keras.Model(inputs, outputs)
9
10 callbacks = [
11     keras.callbacks.ModelCheckpoint("jena_dense.keras",
12                                     save_best_only=True)
13 ]
14 model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
15 history = model.fit(train_dataset,
16                     epochs=10,
17                     validation_data=val_dataset,
18                     callbacks=callbacks)
19
20 model = keras.models.load_model("jena_dense.keras")
21 print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
22 #Visualizing the outcomes

```

```


↗ Epoch 1/10
819/819 [=====] - 50s 60ms/step - loss: 12.4799 - mae: 2.7373 - val_loss: 10.7508 - val_mae: 2.
Epoch 2/10
819/819 [=====] - 41s 50ms/step - loss: 8.4837 - mae: 2.2881 - val_loss: 18.1040 - val_mae: 3.4
Epoch 3/10
819/819 [=====] - 40s 49ms/step - loss: 7.3690 - mae: 2.1315 - val_loss: 10.4222 - val_mae: 2.5
Epoch 4/10
819/819 [=====] - 40s 48ms/step - loss: 6.6881 - mae: 2.0342 - val_loss: 11.1361 - val_mae: 2.6
Epoch 5/10
819/819 [=====] - 50s 60ms/step - loss: 6.1876 - mae: 1.9571 - val_loss: 11.2733 - val_mae: 2.6
Epoch 6/10
819/819 [=====] - 39s 48ms/step - loss: 5.8093 - mae: 1.8973 - val_loss: 10.6586 - val_mae: 2.5
Epoch 7/10
266/819 [=====>.....] - ETA: 19s - loss: 5.5985 - mae: 1.8666

```

```

1 import matplotlib.pyplot as plt
2 loss = history.history["mae"]
3 val_loss = history.history["val_mae"]
4 epochs = range(1, len(loss) + 1)
5 plt.figure()
6 plt.plot(epochs, loss, "bo", label="Training MAE")
7 plt.plot(epochs, val_loss, "b", label="Validation MAE")
8 plt.title("Training and validation MAE")
9 plt.legend()
10 plt.show()
11 #Let's try a 1D convolutional model
12 inputs = keras.Input(shape=(seq_len, original_data.shape[-1]))
13 x = layers.Conv1D(8, 24, activation="relu")(inputs)
14 x = layers.MaxPooling1D(2)(x)
15 x = layers.Conv1D(8, 12, activation="relu")(x)
16 x = layers.MaxPooling1D(2)(x)
17 x = layers.Conv1D(8, 6, activation="relu")(x)
18 x = layers.GlobalAveragePooling1D()(x)
19 outputs = layers.Dense(1)(x)
20 model = keras.Model(inputs, outputs)

```



```

1 callbacks = [
2     keras.callbacks.ModelCheckpoint("jena_conv.keras",
3                                     save_best_only=True)
4 ]
5 model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
6 history = model.fit(train_dataset,
7                     epochs=10,
8                     validation_data=val_dataset,
9                     callbacks=callbacks)
10

```



```

Epoch 1/10
819/819 [=====] - 72s 86ms/step - loss: 22.7385 - mae: 3.7573 - val_loss: 16.7449 - val_mae: 3.
Epoch 2/10
819/819 [=====] - 70s 85ms/step - loss: 15.4664 - mae: 3.1137 - val_loss: 15.1685 - val_mae: 3.
Epoch 3/10
819/819 [=====] - 70s 85ms/step - loss: 13.9005 - mae: 2.9479 - val_loss: 16.0622 - val_mae: 3.
Epoch 4/10
819/819 [=====] - 70s 85ms/step - loss: 13.1099 - mae: 2.8575 - val_loss: 15.2492 - val_mae: 3.
Epoch 5/10
819/819 [=====] - 74s 90ms/step - loss: 12.5349 - mae: 2.7950 - val_loss: 14.8536 - val_mae: 3.
Epoch 6/10
819/819 [=====] - 72s 88ms/step - loss: 12.0707 - mae: 2.7415 - val_loss: 15.7266 - val_mae: 3.
Epoch 7/10
819/819 [=====] - 77s 94ms/step - loss: 11.6463 - mae: 2.6940 - val_loss: 14.4444 - val_mae: 2.
Epoch 8/10
819/819 [=====] - 72s 87ms/step - loss: 11.3081 - mae: 2.6554 - val_loss: 14.8691 - val_mae: 3.
Epoch 9/10
819/819 [=====] - 75s 91ms/step - loss: 11.0393 - mae: 2.6267 - val_loss: 14.8878 - val_mae: 3.
Epoch 10/10
819/819 [=====] - 74s 89ms/step - loss: 10.7928 - mae: 2.5985 - val_loss: 14.2670 - val_mae: 2.

```

```

1 model = keras.models.load_model("jena_conv.keras")
2 print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
3 #Establishing an initial recurrent model as a baseline.
4 #A simple LSTM-based model

```



```

405/405 [=====] - 16s 39ms/step - loss: 15.1933 - mae: 3.0703
Test MAE: 3.07

```

```

1 inputs = keras.Input(shape=(seq_len, original_data.shape[-1]))
2 x = layers.LSTM(16)(inputs)
3 outputs = layers.Dense(1)(x)
4 model = keras.Model(inputs, outputs)

```

```

1 callbacks = [
2     keras.callbacks.ModelCheckpoint("jena_lstm.keras",
3                                     save_best_only=True)
4 ]
5 model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
6 history = model.fit(train_dataset,
7                     epochs=10,
8                     validation_data=val_dataset,
9                     callbacks=callbacks)

```



```

Epoch 1/10
819/819 [=====] - 92s 108ms/step - loss: 40.4353 - mae: 4.6066 - val_loss: 12.2261 - val_mae: 2
Epoch 2/10

```

```

819/819 [=====] - 87s 105ms/step - loss: 10.7716 - mae: 2.5477 - val_loss: 9.9214 - val_mae: 2.
Epoch 3/10
819/819 [=====] - 86s 104ms/step - loss: 9.7521 - mae: 2.4296 - val_loss: 10.1326 - val_mae: 2.
Epoch 4/10
819/819 [=====] - 88s 107ms/step - loss: 9.3576 - mae: 2.3810 - val_loss: 10.5231 - val_mae: 2.
Epoch 5/10
819/819 [=====] - 86s 105ms/step - loss: 9.0289 - mae: 2.3432 - val_loss: 10.0320 - val_mae: 2.
Epoch 6/10
819/819 [=====] - 91s 110ms/step - loss: 8.7527 - mae: 2.3050 - val_loss: 10.2453 - val_mae: 2.
Epoch 7/10
819/819 [=====] - 87s 106ms/step - loss: 8.5766 - mae: 2.2827 - val_loss: 10.4450 - val_mae: 2.
Epoch 8/10
819/819 [=====] - 87s 107ms/step - loss: 8.4186 - mae: 2.2641 - val_loss: 10.3315 - val_mae: 2.
Epoch 9/10
819/819 [=====] - 85s 104ms/step - loss: 8.2340 - mae: 2.2377 - val_loss: 10.5634 - val_mae: 2.
Epoch 10/10
736/819 [=====>....] - ETA: 7s - loss: 8.0807 - mae: 2.2179

```

```

1 model = keras.models.load_model("jena_lstm.keras")
2 print(f"Test MAE: {model.evaluate(test_dataset)[1]:.2f}")
3 #Understanding recurrent neural networks
4 #NumPy implementation of a simple RNN

```

```

↵ 405/405 [=====] - 21s 50ms/step - loss: 11.1081 - mae: 2.6287
Test MAE: 2.63

```

```

1 import numpy as np
2 timesteps = 100
3 input_features = 32
4 output_features = 64
5 inputs = np.random.random((timesteps, input_features))
6 state_t = np.zeros((output_features,))
7 W = np.random.random((output_features, input_features))
8 U = np.random.random((output_features, output_features))
9 b = np.random.random((output_features,))
10 successive_outputs = []
11 for input_t in inputs:
12     output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
13     successive_outputs.append(output_t)
14     state_t = output_t
15 final_output_sequence = np.stack(successive_outputs, axis=0)
16 #A recurrent layer in Keras
17 #An RNN layer that can process sequences of any length

```

```

1 num_features = 14
2 inputs = keras.Input(shape=(None, num_features))
3 outputs = layers.SimpleRNN(16)(inputs)
4 #An RNN layer that returns only its last output step

```

```

1 num_features = 14
2 steps = 120
3 inputs = keras.Input(shape=(steps, num_features))
4 outputs = layers.SimpleRNN(16, return_sequences=False)(inputs)
5 print(outputs.shape)
6 #An RNN layer that returns its full output sequence

```

```

↵ (None, 16)

```

```

1 num_features = 14
2 steps = 120
3 inputs = keras.Input(shape=(steps, num_features))
4 outputs = layers.SimpleRNN(16, return_sequences=True)(inputs)
5 print(outputs.shape)
6 #Stacking RNN layers

```

```

↵ (None, 120, 16)

```

```

1 inputs = keras.Input(shape=(steps, num_features))
2 x = layers.SimpleRNN(16, return_sequences=True)(inputs)
3 x = layers.SimpleRNN(16, return_sequences=True)(x)
4 outputs = layers.SimpleRNN(16)(x)
5 #Exploring advanced techniques with recurrent neural networks
6 #Applying recurrent dropout to combat overfitting
7 #Training and assessing an LSTM model with dropout regularization

```

```

1 inputs = keras.Input(shape=(seq_len, original_data.shape[-1]))
2 x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
3 x = layers.Dropout(0.5)(x)

```



```

1 callbacks = [
2     keras.callbacks.ModelCheckpoint("jena_lstm_dropout.keras",
3                                     save_best_only=True)
4 ]
5 model.compile(optimizer="rmsprop", loss="mse", metrics=["mae"])
6 history = model.fit(train_dataset,
7                     epochs=10,
8                     validation_data=val_dataset,
9                     callbacks=callbacks)
10 inputs = keras.Input(shape=(seq_len, num_features))
11 x = layers.LSTM(32, recurrent_dropout=0.2, unroll=True)(inputs)

Epoch 1/10
819/819 [=====] - 195s 236ms/step - loss: 28.3857 - mae: 3.9216 - val_loss: 9.6649 - val_mae: 2
Epoch 2/10
819/819 [=====] - 188s 229ms/step - loss: 14.7725 - mae: 2.9796 - val_loss: 9.6946 - val_mae: 2
Epoch 3/10
819/819 [=====] - 185s 225ms/step - loss: 13.8027 - mae: 2.8872 - val_loss: 9.1548 - val_mae: 2
Epoch 4/10
819/819 [=====] - 183s 223ms/step - loss: 13.2756 - mae: 2.8263 - val_loss: 8.9818 - val_mae: 2
Epoch 5/10
819/819 [=====] - 186s 227ms/step - loss: 12.7817 - mae: 2.7741 - val_loss: 9.0516 - val_mae: 2
Epoch 6/10
819/819 [=====] - 202s 246ms/step - loss: 12.4087 - mae: 2.7311 - val_loss: 8.9718 - val_mae: 2
Epoch 7/10
819/819 [=====] - 188s 229ms/step - loss: 12.0673 - mae: 2.6935 - val_loss: 9.0521 - val_mae: 2
Epoch 8/10
819/819 [=====] - 186s 227ms/step - loss: 11.6790 - mae: 2.6487 - val_loss: 9.2236 - val_mae: 2
Epoch 9/10
819/819 [=====] - 188s 229ms/step - loss: 11.4827 - mae: 2.6261 - val_loss: 9.8809 - val_mae: 2
Epoch 10/10
819/819 [=====] - 185s 226ms/step - loss: 11.2345 - mae: 2.5966 - val_loss: 9.2737 - val_mae: 2

1 inputs = keras.Input(shape=(seq_len, original_data.shape[-1]))
2 x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs)
3 x = layers.GRU(32, recurrent_dropout=0.5)(x)
4 x = layers.Dropout(0.5)(x)

```